

# vgg-16

March 13, 2025

```
[22]: import numpy as np
      from keras.datasets import cifar10
      from keras.utils import to_categorical
      from keras.applications.vgg16 import VGG16
      from keras.models import Sequential
      from keras.layers import InputLayer, Dense, Flatten, Dropout
      from keras.optimizers import Adam
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      import matplotlib.pyplot as plt

[2]: #load cifar10 dataset
      (x_train, y_train), (x_test, y_test) = cifar10.load_data()

[3]: #normalize the pixel value bw 0 and 1
      x_train = x_train.astype('float32') / 255.0
      x_test = x_test.astype('float32') / 255.0

[4]: #one hot encode the labels (cifar 10 has 10 classes)
      y_train = to_categorical(y_train, 10)
      y_test = to_categorical(y_test, 10)

[5]: #load VGG16 model pre trained on imagenet, without the top layers(fully
      ↪connected layers)
      #base_model = VGG16(weights = 'imagenet' , include_top = False, input_shape =
      ↪(32,32,3))
      base_model = VGG16(weights = 'imagenet' , include_top = False)

[6]: #freeze the layer of VGG16
      for layer in base_model.layers:
          layer.trainable = False

[ ]: #build the model
      model = Sequential()

[8]: #define ip shape for sequential model
      model.add(InputLayer(input_shape=(32,32,3)))
```

```
C:\Users\KH.EN.P2MCA24006\AppData\Local\Packages\PythonSoftwareFoundation.Python
.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-
packages\keras\src\layers\core\input_layer.py:27: UserWarning: Argument
`input_shape` is deprecated. Use `shape` instead.
  warnings.warn(
```

```
[9]: #add the vgg16 base model
model.add(base_model)

[10]: #flatten the op from vgg16's convolutional layers
model.add(Flatten())

[11]: #add a fully connected layer with 256 units and ReLU activaion
model.add(Dense(256, activation='relu'))

[12]: #add a dropout for regularization
model.add(Dropout(0.5))

[13]: #add op layer with 10 units (for 10 classes) and softmax activation
model.add(Dense(10, activation='softmax'))

[14]: #compile the model
model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='categorical_crossentropy', metrics=['accuracy'])

[15]: #print model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14,714,688
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 256)	131,328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2,570

Total params: 14,848,586 (56.64 MB)

Trainable params: 133,898 (523.04 KB)

Non-trainable params: 14,714,688 (56.13 MB)

```
[16]: #data augmentation to reduce overfitting
datagen = ImageDataGenerator(
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
```

```
[17]: #fit the model using data augmentation
batch_size = 64
epochs = 3
train_generator = datagen.flow(x_train,y_train,batch_size=batch_size)
```

```
[18]: #train the model
history = model.fit(train_generator,
                    steps_per_epoch = x_train.shape[0] // batch_size,
                    epochs = epochs,
                    validation_data = (x_test, y_test))
```

Epoch 1/3

```
C:\Users\KH.EN.P2MCA24006\AppData\Local\Packages\PythonSoftwareFoundation.Python
.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-
packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
    self._warn_if_super_not_called()
```

```
781/781          228s 289ms/step -
accuracy: 0.2115 - loss: 2.1797 - val_accuracy: 0.4449 - val_loss: 1.6300
```

Epoch 2/3

```
1/781          3:15 251ms/step -
accuracy: 0.3906 - loss: 1.7655
```

```
C:\Users\KH.EN.P2MCA24006\AppData\Local\Packages\PythonSoftwareFoundation.Python
.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-
packages\keras\src\trainers\epoch_iterator.py:107: UserWarning: Your input ran
out of data; interrupting training. Make sure that your dataset or generator can
generate at least `steps_per_epoch * epochs` batches. You may need to use the
`.repeat()` function when building your dataset.
    self._interrupted_warning()
```


```
781/781          42s 53ms/step -
```

accuracy: 0.3906 - loss: 1.7655 - val\_accuracy: 0.4458 - val\_loss: 1.6294  
Epoch 3/3  
781/781 238s 305ms/step -  
accuracy: 0.4018 - loss: 1.6939 - val\_accuracy: 0.4882 - val\_loss: 1.4868

```
[23]: # Plot Training vs Validation Accuracy
plt.figure(figsize=(14,5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')

plt.savefig('./boo.png')
plt.show()
```



[ ]: